

# DESIGNING MACHINE LEARNING PREDICTION MODEL

Introduction to Artificial Intelligence

Kayode Ohundayo (U2476474)

Year: 2025

# Introduction

This project explores the designing and building of a machine learning prediction model to train customer churn in the banking sector.

Customer churn is defined by the capability of a customer to leave an essential service and here based on my project, it is the likelihood of the customer stopping banking with his bank due to several reasons. Customer churn is important because it helps organisations like banks to take preventive measures for the retention customers.

The project aims to analyse the customer churn dataset thoroughly, preprocess the data, classify and train the model and lastly evaluate the performance of the model.

## Dataset Description

I imported a dataset from Kaggle's customer churn dataset, which contains all the information regarding bank customers and whether they exited the service.

In the project, the data set used will be the Customer Churn data set that holds data about customers of a bank and their churn status.

The data set provides customer demographic data, account data, and activity data. The response variable is **Exited**, which equals one if the customer churned, otherwise zero. Number of rows (10,000)

Our Target variable = **Exited**

Our Type of task = **Classification**

We carried out the data preprocessing on the dataset to prepare it for modeling.

Categorical variables were handled using one-hot encoding. Stratified splitting was used to split the dataset into a training dataset and a testing dataset in a 70:30 ratio.

## Section 1. Data Collection and preprocessing

The project is based on the Bank Customer Churn dataset, which includes 10,000 customers described by demographic, account, and activity information with a churn label **Exited**. I'm solving a binary classification problem: whether a bank customer will churn (**Exited** = 1) or stay (**Exited** = 0).

Main features include:

- **CreditScore**: customer's credit score in numeric format
- **Geography**: customer's country - France, Spain, Germany
- **Gender**: Male/Female

- Age and Tenure (years with bank),
- Balance, NumOfProducts, HasCrCard, IsActiveMember, EstimatedSalary.
- Identity-like fields: RowNumber, Customer Id, Surname

## CODE

```
# =====
# importing the required libraries for the project
# =====

import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

import warnings
warnings.filterwarnings('ignore')
```

```
## load dataset
df = pd.read_csv("Churn_Modelling.csv")
df.head()
df.tail()
df.info()
df.describe()
```

This code loads the churn dataset and provides an overview of feature types, basic statistics, and the presence of the target variable.

**Screenshot solutions // loading the churn modelling dataset and displaying values first 5 rows of the dataset**

```
[89] ✓ 0s
    ## load dataset
    df = pd.read_csv("Churn_Modelling.csv")
    df.head()

...
  RowNumber CustomerId Surname CreditScore Geography Gender Age Tenure Balance NumOfPro
  0         1 15634602 Hargrave        619   France Female  42      2     0.00
  1         2 15647311 Hill           608   Spain Female  41      1  83807.86
  2         3 15619304 Onio            502   France Female  42      8 159660.80
  3         4 15701354 Boni            699   France Female  39      1     0.00
  4         5 15737888 Mitchell        850   Spain Female  43      2 125510.82
```

```
[116] ✓ 0s
    df.tail()

...
  RowNumber CustomerId Surname CreditScore Geography Gender Age Tenure Balance NumOfPro
  9995      9996 15606229 Obijaku        771   France Male   39      5     0.00
  9996      9997 15569892 Johnstone       516   France Male   35     10  57369.61
  9997      9998 15584532 Liu             709   France Female  36      7     0.00
  9998      9999 15682355 Sabbatini       772 Germany Male   42      3  75075.31
  9999     10000 15628319 Walker          792   France Female  28      4 130142.79
```

## Screenshot solutions // displaying the last 5 values rows of the dataset

### 1.2 Missing values and outliers

Data quality was evaluated for missing data, as well as reviewing numeric distributions to determine if there were any observations that could be considered outliers in CreditScore, Balance, or Estimated Salary features, and these outliers would either be accepted as viable or trimmed at extremes as necessary. Missing data was not a factor for any of the columns, so no data had to be imputed.

#### Code

These preprocessing steps assist in ensuring that neither the model is biased by missing data, nor is it dominated by a small number of observations with a particular range of values.

```
## missing values
df.isnull().sum()

import seaborn as sns
import matplotlib.pyplot as plt

# Define list of numeric columns for visualization
```

```

numeric_cols = ['CreditScore', 'Age', 'Tenure', 'Balance',
                'NumOfProducts', 'EstimatedSalary']

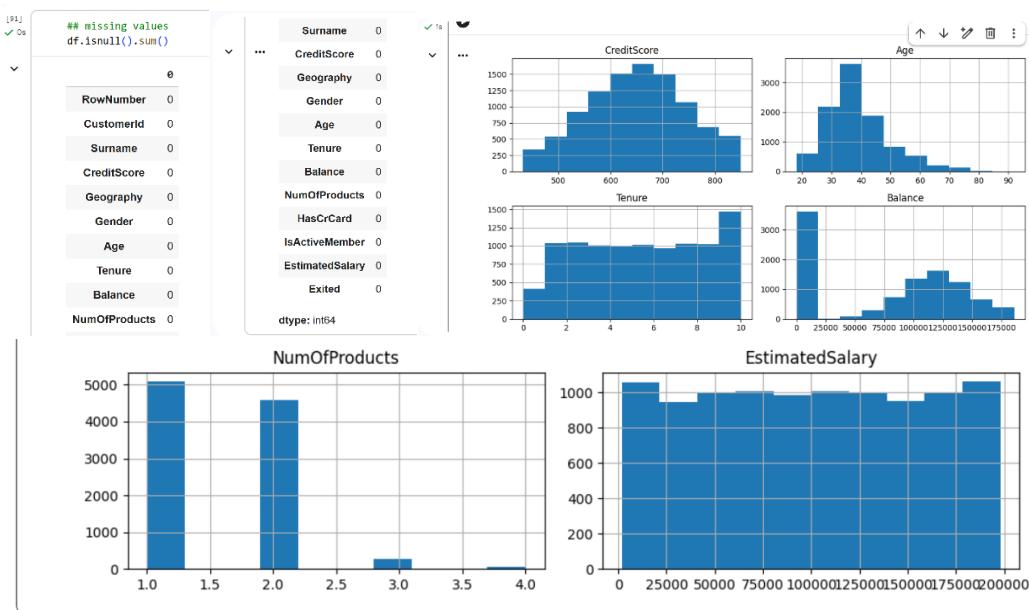
# Create histograms for all numeric features to visualize their distributions
df[numeric_cols].hist(figsize=(10,8))
plt.tight_layout()
plt.show()

# Remove outliers by capping extreme values at 1% and 99% thresholds
for col in ['CreditScore', 'Balance', 'EstimatedSalary']:
    q1 = df[col].quantile(0.01)
    q99 = df[col].quantile(0.99)
    df[col] = df[col].clip(q1, q99)

```

These preprocessing steps help ensure the model is not biased by missing data or dominated by a small number of extreme values, which can distort learning.

## Screenshot solutions



## 1.3 Encoding Categorically Variables

Machine learning models require numerical inputs, so categorical variables were encoded. **Geography** and **Gender** were converted using one-hot (dummy) encoding, with one category dropped to avoid redundancy.

Code

```
# Convert categorical variables to numeric using one-hot encoding
df_encoded = pd.get_dummies(df, columns=['Geography', 'Gender'], drop_first=True)
df_encoded.head()
```

This transformation produces binary features such as **Geography\_Germany**, **Geography\_Spain** and, **Gender\_Male**, making the categorical information usable for both linear and tree-based models.

## Screenshot solutions

The screenshot shows a Jupyter Notebook cell with the following content:

```
[110]  ✓ 0s
▶ # Convert categorical variables to numeric using one-hot encoding
df_encoded = pd.get_dummies(df, columns=['Geography', 'Gender'], drop_first=True)
df_encoded.head()
```

Below the code, the resulting DataFrame is displayed:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	42	2	0.00	1	1	1	101348.88	0
1	608	41	1	83807.86	1	0	1	112542.58	0
2	502	42	8	159660.80	3	1	0	113931.57	0
3	699	39	1	0.00	2	0	0	93826.63	0
4	850	43	2	125510.82	1	1	1	79084.10	0

## Scaling

```
# Split data into features (X) and target variable (y)
X = df_encoded.drop('Exited', axis=1)
y = df_encoded['Exited']

from sklearn.preprocessing import StandardScaler

# Scale numeric features to a standard range
scaler = StandardScaler()
num_cols = ["CreditScore", "Age", "Tenure", "Balance", "EstimatedSalary"]

df_encoded[num_cols] = scaler.fit_transform(df_encoded[num_cols])
```

The dataset was split into training and test sets with a 70/30 ratio, using stratification on **Exited** to preserve the churn proportion. Numeric features were scaled using **StandardScaler** for

models sensitive to feature magnitude (e.g. Logistic Regression), while the Random Forest uses the unscaled version.

### Splitting the data into training and testing sets

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Split data into training (70%) and testing (30%) sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

# Scale features to standardize the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Splitting the data makes it possible to conduct unbiased testing on new data, while scaling helps in improving the speed of convergence of algorithms based on distances and gradients.

## Section 2 – Model selection and Training

### Choice of algorithm and tasks

Because the target **Exited** is binary, both models address a **classification** problem. We chose two models for training:

- Logistic Regression Model: A straightforward linear classifier that yields probabilities and is a good starting point.
- Random Forest Classifier Model: This is a combination of decision trees, which has the ability to identify non-linear relationships. It is common in handling data with columns.

This enables an analysis of a linear versus a non-linear approach on the same problem of churning. **Logistic Regression Model**

```
from sklearn.linear_model import LogisticRegression

# Train a Logistic Regression model
log_reg = LogisticRegression(max_iter=1000, random_state=42)
log_reg.fit(X_train_scaled, y_train)
```

### Screenshot solutions

```
[100] ✓ 0s
from sklearn.linear_model import LogisticRegression

# Train a Logistic Regression model
log_reg = LogisticRegression(max_iter=1000, random_state=42)
log_reg.fit(X_train_scaled, y_train)
```

```
▼          LogisticRegression
LogisticRegression(max_iter=1000, random_state=42)
```

## Random Forest Model

```
from sklearn.ensemble import RandomForestClassifier
# Train a Random Forest model with 200 trees
rf = RandomForestClassifier(
    n_estimators=200,
    max_depth=None,
    random_state=42,
    n_jobs=-1
)
rf.fit(X_train, y_train) # scaling not essential for trees
```

Screenshot solutions

```
[101] ✓ 1s
from sklearn.ensemble import RandomForestClassifier
# Train a Random Forest model with 200 trees
rf = RandomForestClassifier(
    n_estimators=200,
    max_depth=None,
    random_state=42,
    n_jobs=-1
)
rf.fit(X_train, y_train) # scaling not essential for trees
```

```
▼          RandomForestClassifier
RandomForestClassifier(n_estimators=200, n_jobs=-1, random_state=42)
```

## Saving trained models

```
import joblib

# Save trained models and scaler for future use
joblib.dump(log_reg, 'log_reg_model.pkl')
joblib.dump(rf, 'rf_model.pkl')
joblib.dump(scaler, 'scaler.pkl')
```

```
[102] ✓ Os
    import joblib

    # Save trained models and scaler for future use
    joblib.dump(log_reg, 'log_reg_model.pkl')
    joblib.dump(rf, 'rf_model.pkl')
    joblib.dump(scaler, 'scaler.pkl')

    ... ['scaler.pkl']
```

## Section 3 – Prediction and Evaluation

### 3.1 Generating predictions for the test set

Predictions and probabilities were made for both models on the held-out test set.

```
# Make predictions using Logistic Regression model
y_pred_lr = log_reg.predict(X_test_scaled)
y_prob_lr = log_reg.predict_proba(X_test_scaled)[:, 1]

# Make predictions using Random Forest model
y_pred_rf = rf.predict(X_test)
y_prob_rf = rf.predict_proba(X_test)[:, 1]
```

Class labels consist of predicted churn or non churn while probability scores are used for evaluation with threshold measures like ROC and AUC.

### 3.2 Evaluation Metrics

```
from sklearn.metrics import accuracy_score, precision_score, recall_score,
f1_score

# Function to calculate and display model performance metrics
def evaluate_model(y_true, y_pred, name):
    acc = accuracy_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred)
    rec = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
    print(f'{name} - Accuracy: {acc:.3f}, Precision: {prec:.3f}, '
          f'Recall: {rec:.3f}, F1: {f1:.3f}')
    return acc, prec, rec, f1

# Evaluate both models on test data
metrics_lr = evaluate_model(y_test, y_pred_lr, 'Logistic Regression')
metrics_rf = evaluate_model(y_test, y_pred_rf, 'Random Forest')
```

```
[104] ✓ 0s
▶ from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Function to calculate and display model performance metrics
def evaluate_model(y_true, y_pred, name):
    acc = accuracy_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred)
    rec = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
    print(f'{name} - Accuracy: {acc:.3f}, Precision: {prec:.3f}, '
          f'Recall: {rec:.3f}, F1: {f1:.3f}')
    return acc, prec, rec, f1

# Evaluate both models on test data
metrics_lr = evaluate_model(y_test, y_pred_lr, 'Logistic Regression')
metrics_rf = evaluate_model(y_test, y_pred_rf, 'Random Forest')

...
Logistic Regression - Accuracy: 0.813, Precision: 0.628, Recall: 0.196, F1: 0.299
Random Forest - Accuracy: 0.865, Precision: 0.774, Recall: 0.476, F1: 0.590
```

In these metrics accuracy precision recall and f1 measure performance on classification models

These evaluate various aspects of performance: accuracy indicates overall prediction fairness, precision shows reliability of predictions ad recall shows the ability to correctly identify churning customers and f1 is a balance between precision and recall

Model	Accuracy	Precision	Recall	F1
Logistic Regression	0.813	0.628	0.196	0.299
Random Forest	0.865	0.774	0.476	0.590

### 3.3 Comparison and Discussion

For this project, the Random Forest model was overall better at recall and F1 measure than the Logistics Regression model. This means that the Random Forest model is preferable when applied to the churn task. The reason for this preference is that the cost of not contacting a possible chunner (false negative) might be higher than contacting someone who is not going to churn (false positive).

#### 4.1 Confusion matrices

Confusion Matrices were drawn up to illustrate the correct and incorrect results for each model.

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

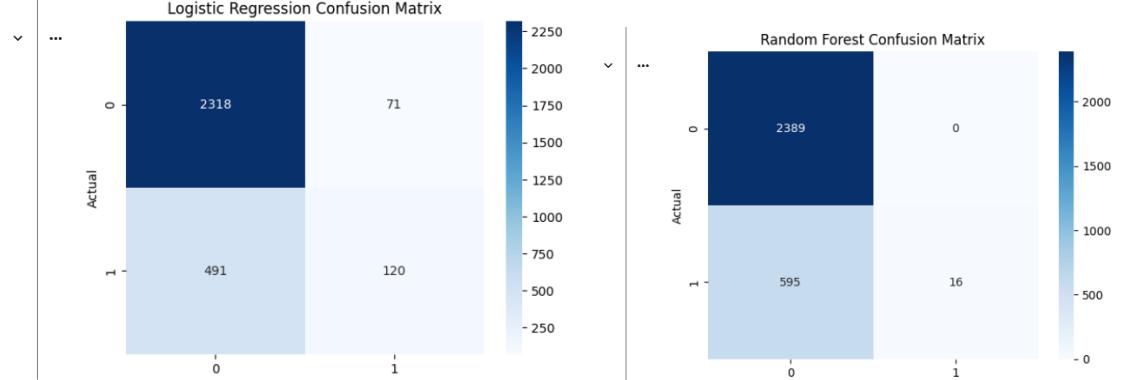
# Function to create and display confusion matrix
def plot_conf_mat(y_true, y_pred, title):
```

```

cm = confusion_matrix(y_true, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(title)
plt.show()

# Plot confusion matrices for both models
plot_conf_mat(y_test, y_pred_lr, 'Logistic Regression Confusion Matrix')
plot_conf_mat(y_test, y_pred_rf, 'Random Forest Confusion Matrix')

```



The plots below emphasize true negatives, false positives, false negatives, and true positives to show more clearly which model yields fewer missed churners:

#### 4.2 ROC curves and AUC

The performance evaluation was done across all classification thresholds with the use of ROC curves. AUC summarized the ranking quality.

```

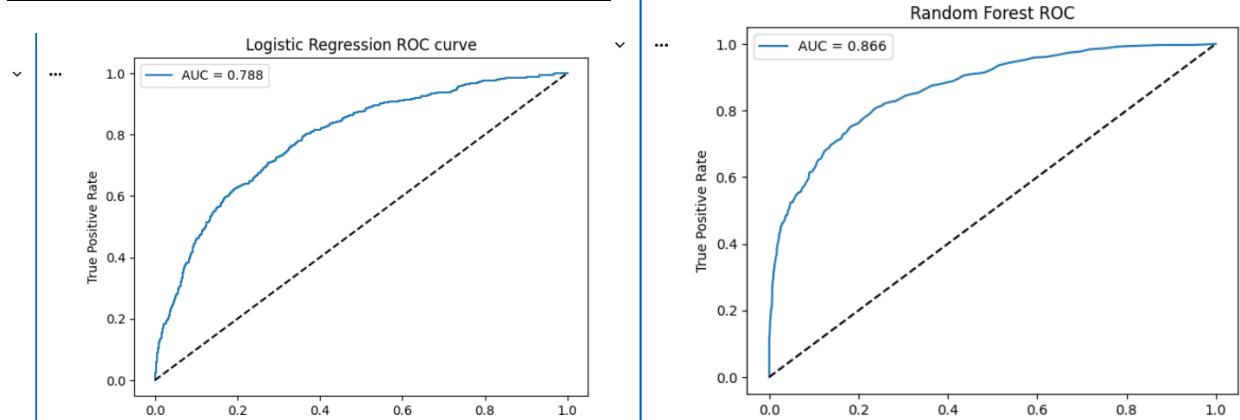
from sklearn.metrics import roc_curve, roc_auc_score

# Function to plot ROC curve and calculate AUC score
def plot_roc(y_true, y_prob, title):
    fpr, tpr, _ = roc_curve(y_true, y_prob)
    auc = roc_auc_score(y_true, y_prob)
    plt.plot(fpr, tpr, label=f'AUC = {auc:.3f}')
    plt.plot([0,1], [0,1], 'k--')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(title)
    plt.legend()
    plt.show()

# Plot ROC curves for both models

```

```
plot_roc(y_test, y_prob_lr, 'Logistic Regression ROC curve')
plot_roc(y_test, y_prob_rf, 'Random Forest ROC')
```



The higher AUC for Random Forest makes it better at ranking churners than Non-Churners, and this makes a strong case for selecting Random Forest as the best model compared to linear regression. The accuracy percentage of random forest is 86% whilst the one for linear regression shows 78%.

#### 4.3 Feature importance

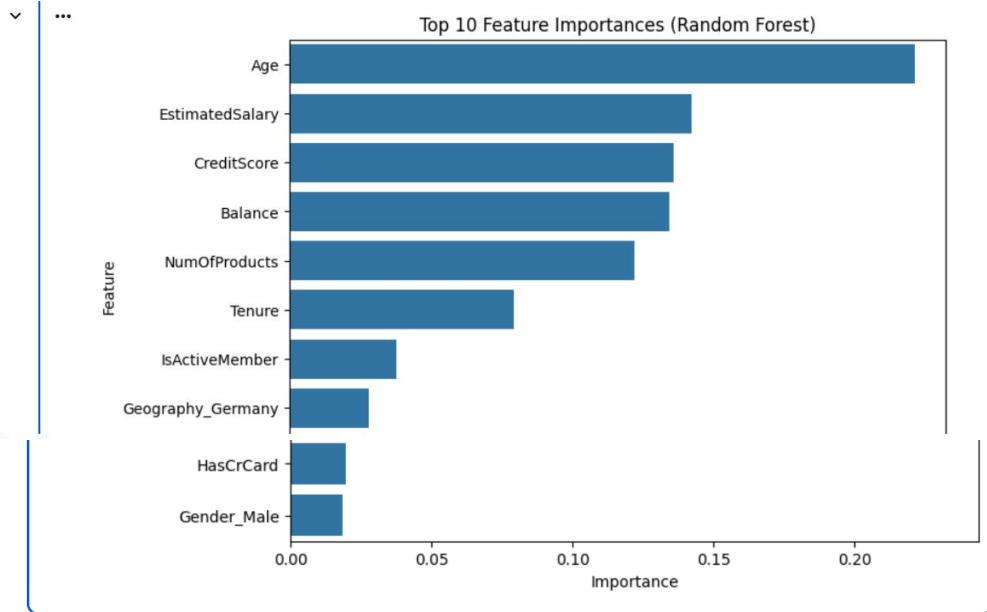
Feature importance obtained from the Random Forest model was also evaluated to identify the most relevant variables or features on which the churn prediction model output relied.

```
import numpy as np

# Extract feature importances from Random Forest model
importances = rf.feature_importances_
feature_names = X.columns

feat_imp = pd.Series(importances,
index=feature_names).sort_values(ascending=False)

# Plot top 10 most important feature
plt.figure(figsize=(8,6))
sns.barplot(x=feat_imp[:10], y=feat_imp.index[:10])
plt.title('Top 10 Feature Importances (Random Forest)')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.show()
```



Main important features from random forest include:

**Age, Balance, NumofProducts, Geography and lastly IsActiveMembers**

This suggest that older customers with certain balance level and lower activity statuses are more likely to churn. This insights can help us to target retention strategies, such as focusing on campaigns on high-risk segments that is identified by the model